

Training curves: validation macro bpb by token-filtering arm
(370M OLMo2, 10B-token corpus)

